

Deep Learning 1




Pointers from Last Time

- Saving your models
 - Just need to save parameters
 - Can Use it anywhere
 - E.g. in linear regression with $y=w_1x_1 + w_2$ you can just save w_1 and w_2 forever!
 - No data dependence after training
- Cost function is only used during training. Prediction function is eternal.
- Linear regression can be made non linear just by adding multiple powers of the same feature

Motivation: ImageNet

Image Classification Dataset



~1,000,000 images
~1,000 classes

Ground truths prepared manually
through Amazon Mechanical Turk

ImageNet Top-5 challenge:

You score if ground truth class is one your top 5 predictions

Best approaches used hand-crafted features
(SIFT, HOGs, Fisher vectors, etc)
+ classifier

Top-5 error rate: ~25%



The Game Has Changed...

Krizhevsky, Sutskever and Hinton;
*ImageNet Classification with
Deep Convolutional Neural networks* **[Krizhevsky12]**

Top-5 error rate of ~15%

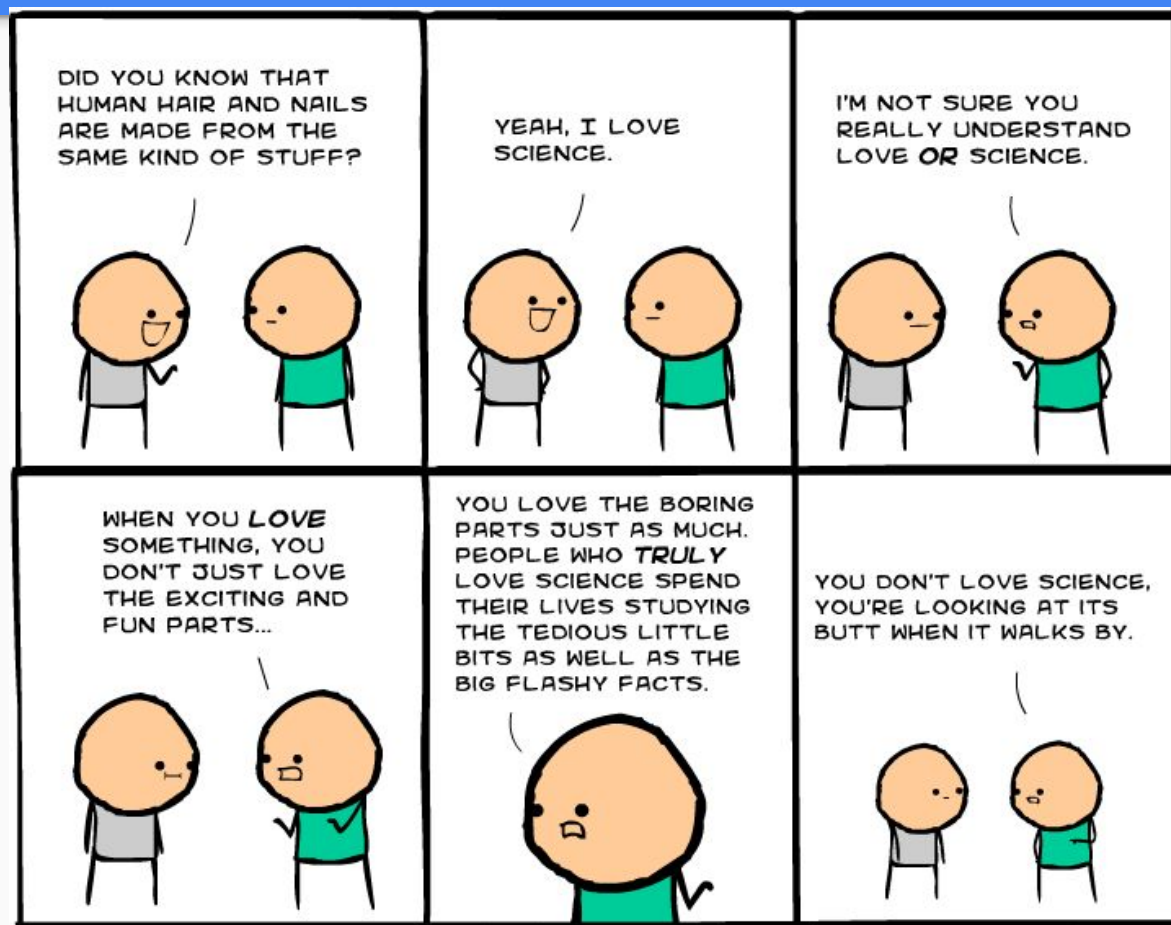
In the last few years, more modern networks
have achieved better results.

[Simonyan14, He15]


Top-5 error rates of ~5-7%

Gradient Descent: Understanding the Math


- We talked about loss functions (i.e. cost functions)
 - Absolute Loss $|h(x) - y|$
 - Quadrative Loss $(h(x) - y)^2$
- Real world is not so generous.
- The following cycle goes on and on until we have a reasonable model
 - Predict using current hypothesis (last time)
 - Find how good the prediction was (last time)
 - **Update the hypothesis. (this time)**
- Q: How Does the update step happen?
 - A: Gradient Descent



Gradient Descent (Cont.)



Gradient Descent (cont.)



Gradient descent algorithm

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$
 (for $j = 1$ and $j = 0$)
}


<https://2.bp.blogspot.com/-AdV-O-MoZHE/TtLibFTaf9I/AAAAAAM/aOxUGP7zl98/s1600/gradient+descent+algorithm+OLS.png>

Repeat until convergence

{
 $\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$
}


<http://2.bp.blogspot.com/-ZxJ87cWjPJ8/TtLtwqv0hCI/AAAAAAAAAV0/9FYqcxJ6dNY/s1600/gradient+descent+algorithm+OLS.png>

Real world is ugly



<https://kousikk.files.wordpress.com/2014/11/screen-shot-2014-11-12-at-2-58-28-pm.png>

TensorFlow



```
import tensorflow as tf
```

$$\begin{matrix} & & \mathbf{W} \\ \mathbf{x} & & \end{matrix}$$

a_0	b_0	c_0
-------	-------	-------

 \cdot

$w_{a,a}$	$w_{a,b}$	$w_{a,c}$
$w_{b,a}$	$w_{b,b}$	$w_{b,c}$
$w_{c,a}$	$w_{c,b}$	$w_{c,c}$

 $=$

a_1	b_1	c_1
-------	-------	-------

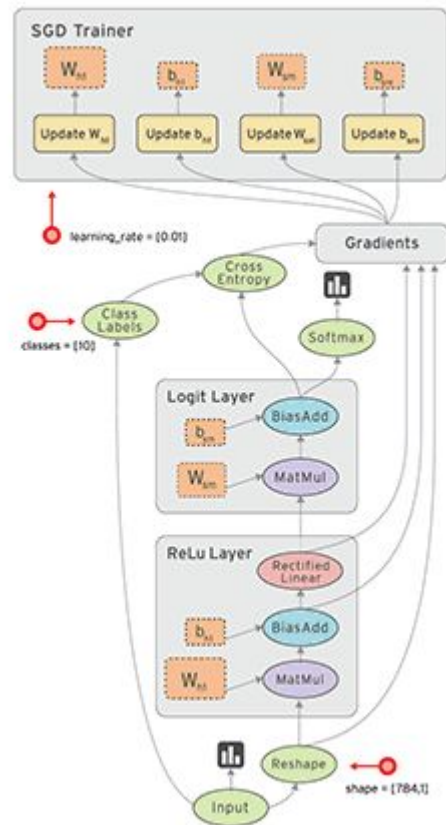
```
y = tf.matmul(x, w)
```

$$\begin{aligned} a_1 &= \text{relu}(a_1) \\ b_1 &= \text{relu}(b_1) \\ c_1 &= \text{relu}(c_1) \end{aligned}$$

```
out = tf.nn.relu(y)
```

Data Flow Graph

- Describe mathematical computation with a directed graph of nodes & edges.
 - Nodes in the graph represent mathematical operations.
 - Edges describe the i/o relationships between nodes.
 - Data edges carry dynamically-sized multidimensional data arrays, or **tensors**.
- Nodes are assigned to computational devices and execute asynchronously and in parallel once all the tensors on their incoming edges becomes available.



TensorFlow Playground

<http://playground.tensorflow.org/>


Back to Code: Hello Tensor

What is Neural Network?


- ❑ Multiple layers
- ❑ Data propagates through layers
- ❑ Transformed by each layer

Possible Reasoning behind effectiveness of NN's

DEEP NETWORKS




A Perceptron: (Logistic Regression)

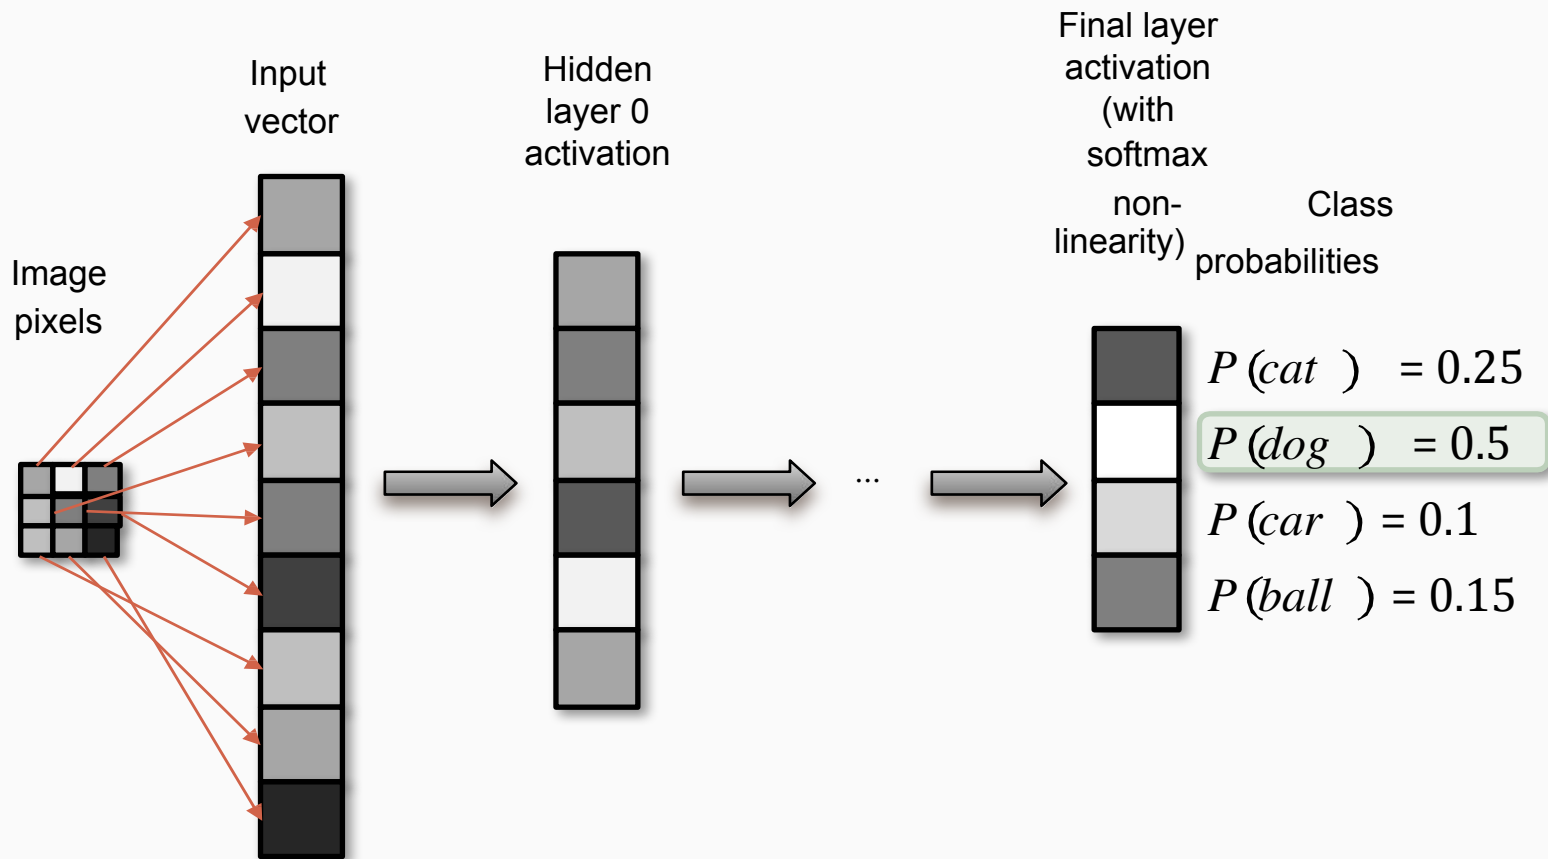


<http://www.parallelr.com/wp-content/uploads/2016/02/neuron.png>

Increasing The classifiers




As a classifier




Non-Linear Activation Function

- Sigmoid: $S(t) = \frac{1}{1 + e^{-t}}$
- Tanh: $\tanh x = \frac{\sinh x}{\cosh x} = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- Rectified Linear Unit (ReLU):
 $f(x) = \max(0, x)$


Most popular activation function for DNN as of 2015, avoids saturation issues, makes learning faster



Sigmoid



Tanh

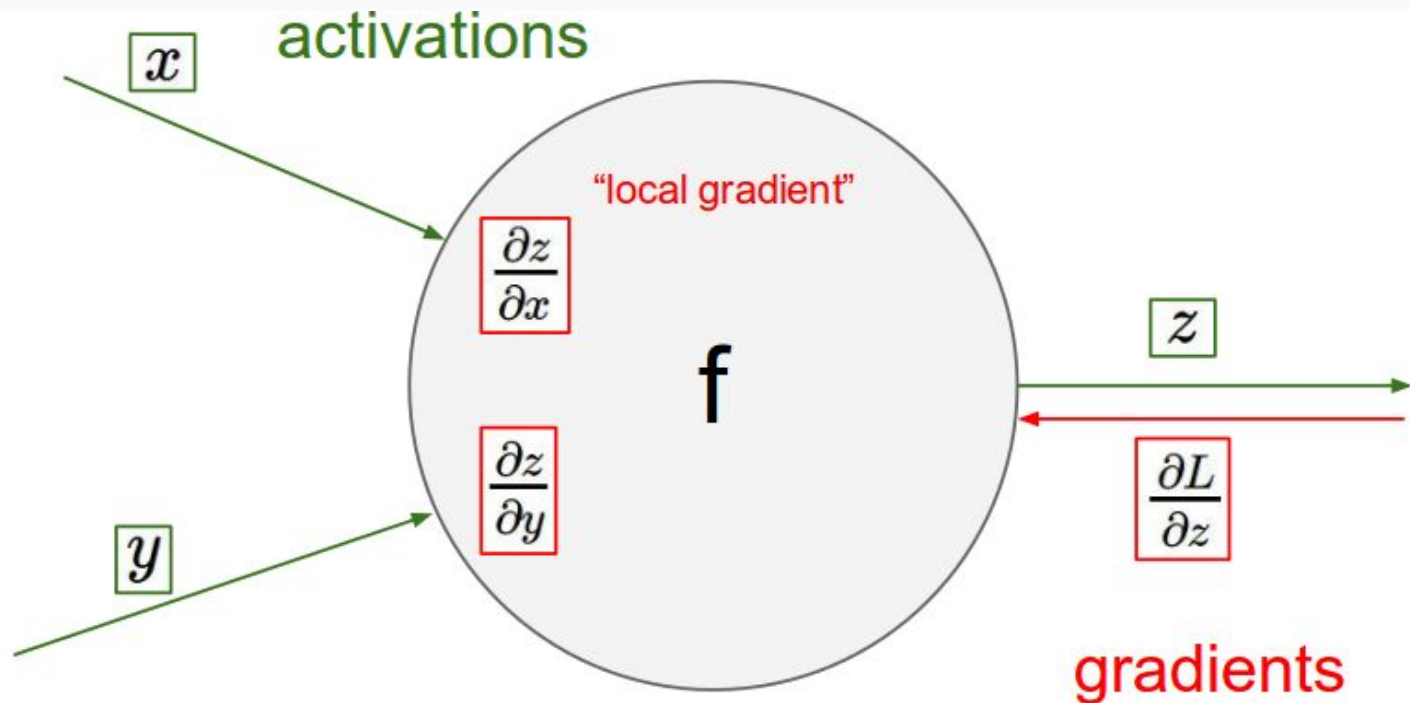


ReLU

Training a Network

- Forward Propagation
 - Multiply Weights by inputs and add them up
 - Pass the result through a nonlinearity (ReLU, Sigmoid)
 - Repeat for next layer until the end layer
- Backward Propagation
 - Each node contributes to a certain degree of error
 - Final errors is a combination of everything
 - Disperse the error to respective nodes
 - Until you reach the starting node

Toy Example Intuition



Toy Example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Toy Example (Cont.)


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$


$$\frac{\partial f}{\partial x}$$

Break (10 Min)

Hands On After That


A Problem

Visualising the learned weights can be educational




Each image visualises the weights connecting pixels to a specific unit in the first hidden layer.

Note the stroke features detected by the various units



The fully connected networks so far have a weakness:

No translation invariance; learned features are
position dependent



How do you solve this?

For more general imagery:

- Requires a training set large enough to see all features in all possible positions...
- Requires network with enough units to represent this...

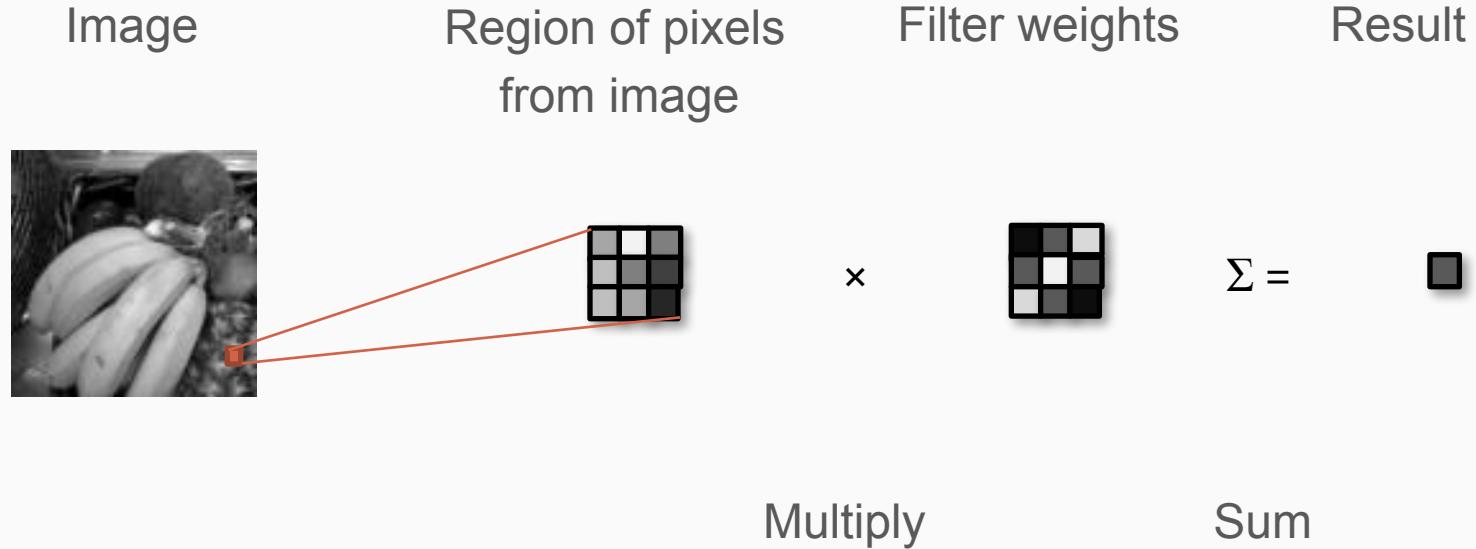
Convolutional Networks

Convolution

Often used for feature detection


Slide a convolutional filter over an image...

Multiply image pixels by filter weights and sum



Do this for all possible positions in the image

Convolution: Gabor filters




Convolution detects features in a position
independent manner



--

Convolutional neural networks learn
position independent filters (feature detectors)

Recap: FC (fully-connected) layer



Convolutional layer

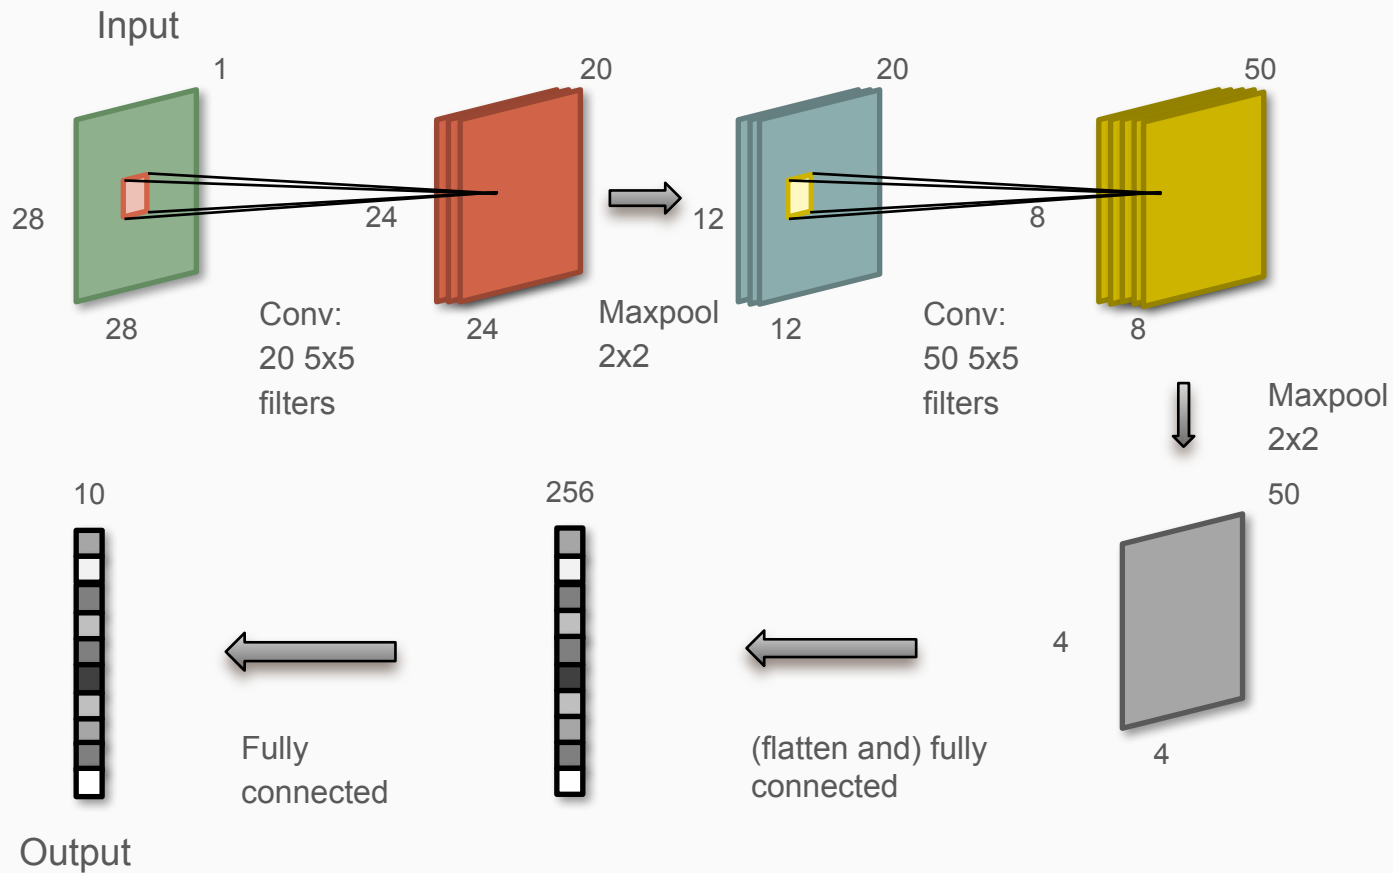


Each unit only connected to units in its neighbourhood

Example:


A Simplified LeNet [**LeCun95**] for MNIST digits

Convolutional layer



Max-pooling 'layer' [Ciresan12]

Take maximum value from each 2×2 pooling region ($p \times p$) in the general case



after 300 iterations over training set:
99.21% validation accuracy

Thanks

Adapted from:

- <https://github.com/Britefury>
- <https://github.com/nlintz/TensorFlow-Tutorials>